

Unsupervised and Reinforcement Learning In Neural Networks Fall 2012


Reinforcement learning

second lecture on RL:


- review from previous week
- on-policy (SARSA) vs off-policy (Q-learning)
- eligibility traces

<http://moodle.epfl.ch/>
Wulfram Gerstner
<http://lcn1.epfl.ch/>

Introduction to reinforcement learning




Theory of Reinforcement learning



○ $s = \text{state}$
 a $a = \text{action}$
 $R_{s \rightarrow s'}^a = \text{reward}$
○ $s' = \text{new state}$

Theory of reinforcement learning – optimal strategy



Optimal strategy:
 - take action a^* with
 $Q(s, a^*) > Q(s, a_i)$
 other actions

But: Q values not known
 (probabilities not known)
 → Estimate Q values

Iterative Update
 $\Delta Q(s, a) = \eta [r - Q(s, a)]$

Reward-based Action Learning

Exploration versus exploitation

Problem: correct Q values not known
 (since reward probabilities are not known)


Exploration versus exploitation

Explore so as to
estimate reward
probabilities

Take action which looks optimal,
so as to maximize reward

Reward-based Action Learning

Exploration versus exploitation



Problem: correct Q values not known
Greedy strategy:
 - take action a^* which looks best
 $Q(s, a^*) > Q(s, a_i)$

ϵ -greedy strategy:
 - take action a^* which looks best
 with prob $P = 1 - \epsilon$

Softmax strategy:
 $P(\text{action } a \mid s) = \frac{\exp[Q(s, a)]}{\sum \exp[Q(s, a')]}$

Optimistic greedy^a:
 start with Q values that are too big

$$Q(s, a) = \sum_{s'} P_{s \rightarrow s'}^a \left[R_{s \rightarrow s'}^a + \gamma \sum_{a'} \pi(s', a') Q(s', a') \right]$$

Bellman equation

$\Delta Q(s, a) = \eta [r - (Q(s, a) + \gamma Q(s', a'))]$

SARSA algo

SARSA algo

Initialise Q values
Start from initial state s

- 1) Being in state s
choose action a
according to policy
- 2) Observe reward r
and next state s'
- 3) Choose action a' in state s'
according to policy
- 4) Update
 $\Delta Q(s, a) = \eta [r - (Q(s, a) + \gamma Q(s', a'))]$
- 5) $s' \rightarrow s$; $a' \rightarrow a$
- 6) Goto 1)

$\Delta Q(s, a) = \eta [r - (Q(s, a) + \gamma Q(s', a'))]$

Reward-based Action Learning

$$Q(s, a) = \sum_{s'} P_{s \rightarrow s'}^a R_{s \rightarrow s'}^a$$

Update

$\Delta Q(s, a) = \eta [r - Q(s, a)]$

$\Delta Q(s, a) = \eta [r - (Q(s, a) + \gamma Q(s', a'))]$

Small learning rate, algos converge in expectation to the correct solution

Reinforcement learning
second lecture on RL:

- ✓ - review from previous week
- ➔ - on-policy (SARSA) vs off-policy (Q-learning)
- eligibility traces

<http://moodle.epfl.ch/>
Wulfram Gerstner
<http://icnl.epfl.ch/>

Exercise from last week: Bellman equation

$$Q(s, a) = \sum_{s'} P_{s \rightarrow s'}^a \left[R_{s \rightarrow s'}^a + \gamma \sum_{a'} \pi(s', a') Q(s', a') \right]$$

Calculated $Q(s, a1)$ and $Q(s, a2)$

a) All actions equal probability
b) Optimal action

$$Q(s, a) = \sum_{s'} P_{s \rightarrow s'}^a \left[R_{s \rightarrow s'}^a + \gamma \sum_{a'} \pi(s', a') Q(s', a') \right]$$

Bellman equations

Resulting Q values are different!

$$Q(s, a) = \sum_{s'} P_{s \rightarrow s'}^a \left[R_{s \rightarrow s'}^a + \gamma \max_{a'} Q(s', a') \right]$$

$\Delta Q(s, a) = \eta [r - (Q(s, a) + \gamma \max_{a'} Q(s', a'))]$

SARSA algo

$\Delta Q(s, a) = \eta [r - (Q(s, a) + \gamma \max_{a'} Q(s', a'))]$

Q-learning algo

Sarsa – on policy temporal difference algo

Initialize $Q(s, a)$ arbitrarily
 Repeat (for each episode):
 Initialize s
 Choose a from s using policy derived from Q (e.g., ϵ -greedy)
 Repeat (for each step of episode):
 Take action a , observe r, s'
 Choose a' from s' using policy derived from Q (e.g., ϵ -greedy)
 $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$
 $s \leftarrow s'; a \leftarrow a'$
 until s is terminal

Action for update
as used for trajectory

Figure 6.9 Sarsa: An on-policy TD control algorithm

$$\Delta Q(s, a) = \eta [r - (Q(s, a) + \gamma Q(s', a'))]$$

From: Reinforcement Learning, Sutton and Barto 1998

Q-learning – off-policy temporal difference algo

Initialize $Q(s, a)$ arbitrarily
 Repeat (for each episode):
 Initialize s
 Repeat (for each step of episode):
 Choose a from s using policy derived from Q (e.g., ϵ -greedy)
 Take action a , observe r, s'
 $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
 $s \leftarrow s'$
 until s is terminal

Policy used for
Action selection

Greedy used for
update

Figure 6.12 Q-learning: An off-policy TD control algorithm.

$$\Delta Q(s, a) = \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

From: Reinforcement Learning, Sutton and Barto 1998

Reinforcement learning

second lecture on RL:

- ✓ - review from previous week
- ✓ - on-policy (SARSA) vs off-policy (Q-learning)
- - eligibility traces

<http://moodle.epfl.ch/>

Wulfam Gerstner

<http://icn1.epfl.ch/>
Exercise from last week

- Update of Q values in SARSA

$$\Delta Q(s, a) = \eta [r - (Q(s, a) + \gamma Q(s', a'))]$$

- policy for action choice:


Pick most often action

$$a_t^* = \arg \max_a Q_a(s, a)$$

Consider a linear sequence of states. Reward only at goal. Actions are up or down.

- Initialise Q values at 0. Start at top. How do Q values develop?
- Q values after 3 complete trials?

goal

**Problem: learning is slow**




- Slow diffusion of information across several states

$$a_t^* = \arg \max_a Q_a(s, a)$$

Consider a linear sequence of states. Reward only at goal. Actions are up or down.

- Initialise Q values at 0. Start at top. How do Q values develop?
- Q values after 3 complete trials?

goal

**Eligibility trace**Sarsa $\Delta Q(s, a) = \eta [r - (Q(s, a) + \gamma Q(s', a'))]$ 

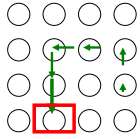
Idea: at the moment of reward update also previous action values along trajectories

Eligibility trace - algo

$$\text{Sarsa}(0) \quad \Delta Q(s,a) = \eta [r + \gamma Q(s',a') - Q(s,a)]$$

$$\delta_t = r_{t+1} - [Q_a(s,a) - \gamma \cdot Q_a(s',a')]$$

$$\text{Sarsa}(\lambda) \quad \Delta Q_{aj} = \eta \delta_t e_{aj}$$


 if $a = \text{action taken in state } j$
 $e_{aj}(t) = \gamma \lambda e_{aj}(t - \Delta t) + 1$
 else
 $e_{aj}(t) = \gamma \lambda e_{aj}(t - \Delta t)$

Eligibility trace – Sarsa()

```

Initialize  $Q(s,a)$  arbitrarily
Repeat (for each episode):
  Initialize  $s, a$  and  $e(s,a)=0$ 
  Repeat (for each step of episode):
    Take action  $a$ , observe  $r, s'$ 
    Choose  $a'$  from  $s'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
     $\delta \leftarrow r + \gamma Q(s',a') - Q(s,a)$ 
     $e(s,a) \leftarrow e(s,a) + 1$ 
    For all  $s, a$ :
       $Q(s,a) \leftarrow Q(s,a) + \alpha \delta e(s,a)$ 
       $e(s,a) \leftarrow \gamma \lambda e(s,a)$ 
     $s \leftarrow s'; a \leftarrow a'$ 
  until  $s$  is terminal
    
```

Figure 7.11 Tabular Sarsa(λ).

From: Reinforcement Learning
Sutton and Barto 1998

Reinforcement learning

second lecture on RL:



- ✓ - review from previous week
- ✓ - on-policy (SARSA) vs off-policy (Q-learning)
- ✓ - eligibility traces
- - continuous state space

<http://moodle.epfl.ch/>

Wulfram Gerstner


<http://icn1.epfl.ch/>

Continuous states – example



Continuous states – example

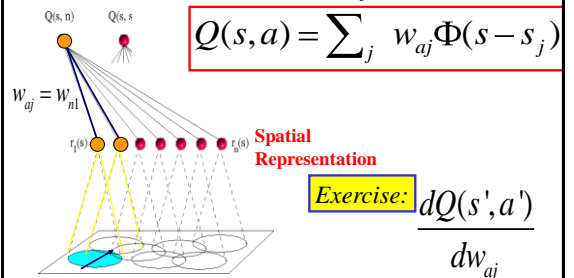
Blackboard:



Continuous states

$$Q(s,a) = \sum_j w_{aj} r_j(s)$$

$$Q(s,a) = \sum_j w_{aj} \Phi(s - s_j)$$



Bellman equation

$$Q(s, a) = \sum_{s'} P_{s \rightarrow s'}^a \left[R_{s \rightarrow s'}^a + \gamma \sum_{a'} \pi(s', a') Q(s', a') \right]$$

Blackboard: idea

$$Q(s, a) = r + \gamma Q(s', a')$$

Eligibility and continuous: TD()

- TD error in SARSA**

$$\delta_t = R_{t+1} - Q_a(s, a) - \gamma \cdot Q_{a'}(s', a')$$
- Function approximation**

$$Q_w(s, a) = \sum_{i=1}^n w_i^a \cdot r_i(s)$$
- policy for action choice:**
Pick most often action

$$a_t^n = \arg \max_a Q_a(s, a)$$
- Eligibility trace**

$$\Delta w_{aj} = \eta \delta_t e_{aj}$$

$$e_{aj}(t) = \gamma \lambda e_{aj}(t - \Delta t) + \begin{cases} r_j & \text{if } a = \text{action taken} \\ 0 & \end{cases}$$

memory

Eligibility and continuous: TD()

- TD error in SARSA**

$$\delta_t = R_{t+1} - Q_a(s, a) - \gamma \cdot Q_{a'}(s', a')$$
- Function approximation**

$$Q_w(s, a) = \sum_{i=1}^n w_i^a \cdot r_i(s)$$
- Synaptic update**

$$\Delta w_{aj} = \eta \delta_t e_{aj}$$
- Eligibility trace (memory at synapse):**

$$e_{aj}(t) = \gamma \lambda e_{aj}(t - \Delta t) + \begin{cases} r_j \cdot 1 & \text{if } a = a' (= \text{action taken}) \\ r_j \cdot 0 & \text{else} \end{cases}$$

Biological interpretation: next week

Reward-based Action Learning

Connection reinforced if action a at state s successful

Success signal

Learning Rule

$$\Delta w_{aj} = \eta \delta_t e_{aj}$$

Spatial Representation

Example – mountain car

214 Generalization and Function Approximation

reward $r = -1$ except at goal

Goal

Step 425

Episode 12

Episode 104

Episode 1000

Episode 9000

Figure 8.10 The mountain-car task (upper left panel) and the cost-to-go function ($-\max_a Q_a(s, a)$) learned during one run.


Consider the task of driving an underpowered

Biological Principles of Learning: spatial learning

Validating the Model

The KHEPERA mobile miniature robot

Experimental arena





422 x 316 pixels



Open-field Navigation Experiments

(Biol. Cybern., 2000)



Navigation map
after 20 training trials



11.3 The Acrobot

goal: Raise tip above line




Figure 11.4 The acrobot.

Steps per episode
(log scale)




Figure 11.6 Learning curves for Sarsa(λ) on the acrobot task.

274

Case Studies




Figure 11.7 A typical learned behavior of the acrobot. Each group is a series of consecutive positions, the thicker line being the first. The arrow indicates the torque applied at the second joint.

The end